

# Branch-and-Price for the Traveling Tournament Problem (TTP):

## Implementation Notes for `BNP_NLX.py`

September 28, 2025

### 1 Problem statement and notation

Let the set of teams be  $T = \{1, \dots, n\}$  and the set of time slots for a double round-robin be  $S = \{1, \dots, 2(n-1)\}$ . Distances are given by a matrix  $D = (d_{ij})_{i,j \in T}$ . Each team plays every other team twice (home/away), subject to *break constraints* limiting consecutive home/away streaks to the interval  $[L, U]$ . Optional *non-repeater constraints* (NRCs) forbid immediate home-and-away reversals across consecutive slots.

A *tour* for team  $t$  is a feasible assignment of opponents and home/away decisions over all slots. In the branch-and-price (B&P) scheme, tours are the *columns*.

### 2 Dantzig–Wolfe decomposition and the the RMP

Let  $\lambda_{t,p} \in [0, 1]$  denote whether tour  $p$  (from the current pool) is selected for team  $t \in T$ . The Restricted Master Problem (RMP) is a set-partitioning model:

$$\min_{\lambda} \sum_{t \in T} \sum_{p \in \mathcal{P}_t} c_{t,p} \lambda_{t,p} \quad (1)$$

$$\text{s.t.} \quad \sum_{p \in \mathcal{P}_t} \lambda_{t,p} = 1, \quad \forall t \in T \quad (\text{one tour per team}) \quad (2)$$

$$\sum_{p \in \mathcal{P}_t} a_{t,p,s} \lambda_{t,p} = 1, \quad \forall (t, s) \in T \times S \quad (\text{slot coupling}) \quad (3)$$

$$\sum_{t \in T} \sum_{p \in \mathcal{P}_t} b_{t,p,i,s,j} \lambda_{t,p} \leq 1, \quad \forall (i, s, j) \quad (\text{NRCs, separated}) \quad (4)$$

$$0 \leq \lambda_{t,p} \leq 1 \quad (\text{LP; IP solves for UB}).$$

Here  $a_{t,p,s} \in \{0, 1\}$  indicates that tour  $p$  for team  $t$  plays in slot  $s$ , and  $b_{t,p,i,s,j} \in \{0, 1\}$  aggregates the NRC pattern “ $i$  hosts  $j$  at  $s$  and  $j$  hosts  $i$  at  $s+1$ ” on a tour  $p$ .

**Code map.** RMP construction/solve is handled by `MasterLP`. Dynamic separation of NRCs (4) is implemented by `separate_nrc_cuts`. Columns are inserted when pricing returns negative reduced-cost tours.

### 3 Duals and reduced costs

Let  $\pi_t$  be the dual of (2) for team  $t$ , let  $\mu_{t,s}$  be the dual of (3) for  $(t, s)$ , and let  $\beta_{i,s,j}$  be the dual of an active NRC inequality for  $(i, s, j)$ . The reduced cost of a tour  $p \in \mathcal{P}_t$  is

$$\tilde{c}_{t,p} = c_{t,p} - \pi_t - \sum_{s \in S} \mu_{t,s} y_{t,p}(s) - \sum_{(i,s,j)} \beta_{i,s,j} z_{t,p}(i, s, j), \quad (5)$$

where  $y_{t,p}(s) \in \{0, 1\}$  captures the slot usage of  $p$ , and  $z_{t,p}(i, s, j) \in \{0, 1\}$  captures NRC patterns induced by  $p$ . A negative  $\tilde{c}_{t,p}$  indicates a profitable column to add.

**Arc-based accounting.** Rather than compute (5) at the tour level, the code pushes dual terms onto *arcs* of a time-expanded network for team  $t$  and solves a shortest path with resource constraints (SPPRC). For an arc used at slot  $s$  from state  $u$  to  $v$  with travel  $d_{uv}$ , the arc’s reduced cost is

$$\tilde{c}_{t,s}^{\text{arc}}(u \rightarrow v) = d_{uv} - \mu_{t,s} - (\beta\text{-term if an NRC pattern is triggered}). \quad (6)$$

The implementation uses integer-scaled duals (e.g., `SCALE=106` via `to_int_dual`) to avoid numerical tolerance issues.

**Code map.** Arc costs with  $(\mu, \beta)$  are built in `build_rc_edges_with_beta_int`; forward/backward DP potentials in `rc_forward_dp_int` and `rc_backward_dp_int`; labeling-based pricing in `pricing_single_best_`.

### 4 Pricing network and resource feasibility

The time-expanded network for a fixed team  $t$  includes nodes encoding slot index, venue/opponent, break-streak state, and a bitmask of visited opponents. The SPPRC enforces:

- **Elementarity:** each opponent appears at most once from  $t$ ’s away perspective (bitmask resource).
- **Break limits:**  $L \leq$  consecutive home/away streak  $\leq U$ ; streak resource is updated along arcs.
- **Branching filters:** forced home/away and include/exclude specific games are implemented by removing or forcing arcs in the relevant slots.
- **Symmetry reduction:** at the middle boundary between halves (slot  $n-1$ ), away opponents are restricted (e.g., “opponent index  $< t$ ”), removing mirror-symmetric solutions; see the `symmetry_middle` option in `build_rc_edges_with_beta_int`.

### 5 Dynamic NRC separation

After each RMP solve, the fractional  $\lambda$ -solution is projected to slot-level quantities to form  $y(i, s, j)$  values. Any violated NRC of type

$$y(i, s, j) + y(j, s+1, i) \leq 1 \quad (7)$$

is added to the RMP, and the LP is re-solved. Duals  $\beta$  from these cuts then feed into arc reduced costs via (6).

Code map. `separate_nrc_cuts`.

## 6 Exact pruning in pricing via bidirectional DP

Let  $fw[(s, u)]$  denote the best reduced cost from the source to node  $(s, u)$ , and  $bw[(s, u)]$  the best reduced cost from  $(s, u)$  to the sink, computed on the *arc-reduced-cost* network ignoring resource conflicts. For an arc  $(s, u \rightarrow v)$  with cost  $r = \tilde{c}_{t,s}^{\text{arc}}(u \rightarrow v)$ , if

$$fw[(s, u)] + r + bw[(s+1, v)] - \pi_t \geq 0, \quad (8)$$

then no completion through this arc can produce a negative reduced-cost tour and the arc may be safely eliminated *before* labeling. The code computes integer  $fw, bw$  via `rc_forward_dp_int/rc_backward_dp_int` and prunes arcs accordingly; the remaining graph is then searched by labeling.

## 7 Column generation loop

The overall CG loop (see `column_generation_with_branch`) is:

1. Solve the RMP (LP) and extract duals  $(\pi, \mu, \beta)$ .
2. For each team  $t \in T$ , run `pricing_single_best_int` to search for a tour with  $\tilde{c}_{t,p} < 0$ .
3. If any negative columns are found, add them to the RMP and return to step 1.
4. If none are found, solve the RMP as an IP to get a feasible schedule (upper bound, UB), and then apply reduced-cost fixing (Proposition A) to shrink the RMP.

## 8 Reduced-cost fixing (Proposition A)

Let LB be the current RMP value and UB the best incumbent. Any column with integer reduced cost  $\tilde{c} \geq (\text{UB} - \text{LB})$  can be fixed at  $\lambda = 0$  without losing optimality. The implementation computes integer  $\tilde{c}$  for active columns and tightens variable bounds; see `apply_proposition_A`.

## 9 Branching and strong branching

If the LP solution remains fractional after CG and cut separation, the algorithm performs best-bound B&P search (see `branch_and_price_bestbound`). The branching rules are compatible with pricing:

- **Include/Exclude event:** force or forbid a specific game (home  $i$  vs  $j$  at slot  $s$ ).
- **Force H/A at slot:** for a given  $(t, s)$ , force home or away.

These are enforced in the RMP by equalities/zero-sums over  $\lambda$  columns and respected in pricing by filtering arcs in the affected slots (see `ensure_branch_feasibility` and `build_rc_edges_with_beta_int`). Candidate selection uses event strong branching (`rank_fractional_events`) with fallback to home/away (`pick_fractional_HA`).

## 10 Initialization (seeding) and symmetry

To start CG with a feasible RMP, the code builds canonical round-robin schedules (Berger tables; `generate_single_rr`, `generate_double_rr`), constructs team tours obeying break limits (`build_seed_tours`), and applies light local improvements (2-opt within halves; `improve_team_tour_by_2opt`). The option `symmetry_middle` restricts the opponent set at the half boundary to kill mirrored solutions early.

## 11 Bounds, search, and termination

At each node, the LP objective provides a lower bound (LB). Solving the RMP integrally yields feasible schedules and updates the global UB. Nodes with  $LB \geq UB$  are fathomed. The search terminates when the open-node list is empty or when  $LB = UB$ . Final schedules and travel statistics are reported via `print_solution`.

## 12 Compact mathematical formulation (self-contained)

For a *compact*, cross-reference-free summary (avoids ?? if you compile once), we give a re-stated master and pricing with local tags:

$$\min_{\lambda} \sum_{t \in T} \sum_{p \in \mathcal{P}_t} c_{t,p} \lambda_{t,p} \quad (\text{M0})$$

$$\text{s.t.} \quad \sum_{p \in \mathcal{P}_t} \lambda_{t,p} = 1, \quad \forall t \in T \quad (\text{one tour}) \quad (\text{M1})$$

$$\sum_{p \in \mathcal{P}_t} a_{t,p,s} \lambda_{t,p} = 1, \quad \forall (t, s) \in T \times S \quad (\text{coupling}) \quad (\text{M2})$$

$$\sum_{t \in T} \sum_{p \in \mathcal{P}_t} b_{t,p,i,s,j} \lambda_{t,p} \leq 1, \quad \forall (i, s, j) \quad (\text{NRC}) \quad (\text{M3})$$

$$0 \leq \lambda_{t,p} \leq 1. \quad (\text{M4})$$

Tour reduced costs and arc reduced costs for pricing (team  $t$ ) are:

$$\tilde{c}_{t,p} = c_{t,p} - \pi_t - \sum_{s \in S} \mu_{t,s} y_{t,p}(s) - \sum_{(i,s,j)} \beta_{i,s,j} z_{t,p}(i, s, j), \quad (\text{P1})$$

$$\tilde{c}_{t,s}^{\text{arc}}(u \rightarrow v) = d_{uv} - \mu_{t,s} - (\beta\text{-term if NRC is triggered}). \quad (\text{P2})$$

Bidirectional DP pruning condition:

$$fw[(s, u)] + \tilde{c}_{t,s}^{\text{arc}}(u \rightarrow v) + bw[(s+1, v)] - \pi_t \geq 0 \implies \text{prune arc } (s, u \rightarrow v). \quad (\text{P3})$$

Reduced-cost fixing threshold (Proposition A): fix columns with  $\tilde{c} \geq (UB - LB)$ .